

Recurrent Neural Networks

Remembering Past Data

Dr. Mauricio Toledo-Acosta

Diplomado Ciencia de Datos con Python

Table of Contents

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

1 Recurrent Neural Networks

2 LSTM

Recurrent Neural Networks

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- **Recurrent neural networks are designed for sequential data** like text sentences or time-series.

Recurrent Neural Networks

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- **Recurrent neural networks are designed for sequential data** like text sentences or time-series.
- The input is of the form x_1, \dots, x_n , where x_t is a d -dimensional point received at the time-stamp t .

Recurrent Neural Networks

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- **Recurrent neural networks are designed for sequential data** like text sentences or time-series.
- The input is of the form x_1, \dots, x_n , where x_t is a d -dimensional point received at the time-stamp t .
- **Successive data are dependent on past data.** Therefore, it is helpful to receive a particular input x_t only after the earlier inputs have already been received and converted into a hidden state.

Recurrent Neural Networks

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- **Recurrent neural networks are designed for sequential data** like text sentences or time-series.
- The input is of the form x_1, \dots, x_n , where x_t is a d -dimensional point received at the time-stamp t .
- **Successive data are dependent on past data.** Therefore, it is helpful to receive a particular input x_t only after the earlier inputs have already been received and converted into a hidden state.
- The traditional type of feed-forward network in which all inputs feed into the first layer does not achieve this goal.

Recurrent Neural Networks

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.

Recurrent Neural Networks

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below

Recurrent Neural Networks

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input x_t at each time-stamp, and a hidden state h_t that changes at each time stamp as new data points arrive.

Recurrent Neural Networks

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input x_t at each time-stamp, and a hidden state h_t that changes at each time stamp as new data points arrive.
- Each time-stamp also has an output value y_t . For example,

Recurrent Neural Networks

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input x_t at each time-stamp, and a hidden state h_t that changes at each time stamp as new data points arrive.
- Each time-stamp also has an output value y_t . For example,
 - In a time-series setting, the output y_t is the forecasted prediction of x_{t+1} .

Recurrent Neural Networks

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input x_t at each time-stamp, and a hidden state h_t that changes at each time stamp as new data points arrive.
- Each time-stamp also has an output value y_t . For example,
 - In a time-series setting, the output y_t is the forecasted prediction of x_{t+1} .
 - In the text-setting, we are predicting the next word.

Recurrent Neural Networks

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input x_t at each time-stamp, and a hidden state h_t that changes at each time stamp as new data points arrive.
- Each time-stamp also has an output value y_t . For example,
 - In a time-series setting, the output y_t is the forecasted prediction of x_{t+1} .
 - In the text-setting, we are predicting the next word.
 - In some applications, we do not output y_t at each time stamp, but only at the end of the sequence (sentiment analysis).

Examples of RNN

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- Fully recurrent neural networks (FRNN)

Examples of RNN

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- Fully recurrent neural networks (FRNN)
- Independently RNN (IndRNN)

Examples of RNN

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- Fully recurrent neural networks (FRNN)
- Independently RNN (IndRNN)
- Long short-term memory (LSTM)

Examples of RNN

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM


- Fully recurrent neural networks (FRNN)
- Independently RNN (IndRNN)
- Long short-term memory (LSTM)
- Gated recurrent units (GRUs)

General Architecture of a RNN

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM



The vanishing gradient problem

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.

The vanishing gradient problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.
- In these methods, during each iteration of training each of the neural network's weights receives an update proportional to the partial derivative of the error function.

The vanishing gradient problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.
- In these methods, during each iteration of training each of the neural network's weights receives an update proportional to the partial derivative of the error function.
- In some cases, the gradient will be very small, preventing the weight from changing its value.

The vanishing gradient problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.
- In these methods, during each iteration of training each of the neural network's weights receives an update proportional to the partial derivative of the error function.
- In some cases, the gradient will be very small, preventing the weight from changing its value.
- For example, the hyperbolic tangent function have gradients in the range $(0, 1]$, therefore, the gradient decreases exponentially with the number of layers.

The vanishing gradient problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.
- In these methods, during each iteration of training each of the neural network's weights receives an update proportional to the partial derivative of the error function.
- In some cases, the gradient will be very small, preventing the weight from changing its value.
- For example, the hyperbolic tangent function have gradients in the range $(0, 1]$, therefore, the gradient decreases exponentially with the number of layers.
- With activation functions whose derivatives can take on larger values, one risks encountering the related opposite exploding gradient problem.

Table of Contents

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

1 Recurrent Neural Networks


2 LSTM

Long short-term memory (LSTM)

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM



LSTM

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- Unlike standard feedforward neural networks, **LSTM has feedback connections**. LSTMs can process not only single data points, but also entire sequences of data.

LSTM

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- Unlike standard feedforward neural networks, **LSTM has feedback connections**. LSTMs can process not only single data points, but also entire sequences of data.
- **A LSTM unit is composed of a cell, an input gate, an output gate and a forget gate**. The cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell.

LSTM

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

- Unlike standard feedforward neural networks, **LSTM has feedback connections**. LSTMs can process not only single data points, but also entire sequences of data.
- **A LSTM unit is composed of a cell, an input gate, an output gate and a forget gate**. The cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell.
- LSTM networks are well-suited to classifying, and making predictions based on time series data.

LSTM

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM


- Unlike standard feedforward neural networks, **LSTM has feedback connections**. LSTMs can process not only single data points, but also entire sequences of data.
- **A LSTM unit is composed of a cell, an input gate, an output gate and a forget gate**. The cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell.
- LSTM networks are well-suited to classifying, and making predictions based on time series data.
- **LSTMs were developed to deal with the vanishing gradient problem** that can be encountered when training traditional RNNs.

LSTM

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM



c_t : Cell state, *memoría de la célula*

f_t : Forget gate, *¿qué tanto olvidamos de c_t*

o_t : Output gate

i_t : Input gate

$i_t * \tilde{C}_t$: *La contribución de esta entrada a la memoria*

h_t : Hidden-state vector, output

LSTM: Unfolded

Recurrent
Neural
Networks

Recurrent
Neural
Networks

LSTM

